

# Formulaire - MAC

Leonard Cseres | November 12, 2025

**Impedance Mismatch Objet-Relationnel** Disconnect between data layer and application layer

- Solutions: *ORM, NoSQL*

## NoSQL

- Horizontal scaling
- No schema, fields can be added later
- Easy replication
- Simple API
- Not ACID

## Column oriented

- Écriture à grande échelle
- Accès aux données co-localisé (pour la lecture et l'écriture)

## Key-value

- Caches
- Domaine simple avec accès en lecture rapide
- Systèmes massivement concurrents
- Opaque value

## Document oriented

- Suited for agile dev.
- When data modeling follows the structures of natural documents
- No need for migration schema
- No need for ORM layer
- Replicated, each document is independent
- Separation approach: normalized data
  - + Data consistency through a single canonical source
  - + Simpler queries (closer to relational databases)
  - + Better cache efficiency and hardware utilization
  - Requires multiple lookups and joins
  - Forced consistency may be undesirable in some contexts
- Imbrication approach: unnormalized data
  - + Faster access (no joins, single document retrieval)
  - + Fewer failure points in distributed systems
  - + Simpler application logic
  - Risk of data inconsistency due to redundancy
  - More complex queries on nested data
  - Larger, heavier documents

## Models

- Represent n-to-m relations
  - document model: *difficult*
  - relation model: *easy*
- Optional fields
  - document model: *possible*
  - relations model: *not possible*

## Couchbase

Clause	Scope of alias
WITH	Anywhere
FROM	Anywhere
LET	Anywhere

Clause	Scope of alias
LETTING	HAVING, SELECT and ORDER BY clauses
SELECT	SELECT and ORDER BY clauses
FOR	The local collection expression

## Indexes

- Primary (on the document key): `CREATE PRIMARY INDEX ... ON ...;`
- Secondary (on any key-value): `CREATE INDEX ... ON ...;`
- Secondary composite (on any key-value): `CREATE INDEX ... ON ...();`
- Covering (the query only selects filed in the index, no need to go to the document)

## Graph

- Interconnected data
- When the domain can be represented by nodes and relations
- Social media, recommendation engines
- Relational DBs compute the relations during the query; graph DBs store them.
- Native: custom underlying storage
- Non-Native: underlying relational DB
- Index-free adjacency:** the relations are stored instead of being indexed

**Why distribute data?** Scalability (Evolutivité), fault tolerance/high availability, reduced latency

## Shared memory architecture

- Vertical scaling (scale up)
- Processors share unique memory (SMP or symmetric multiprocessing)
  - higher costs
  - limited data growth
  - limited fault tolerance

## Shared disks architecture

- Independent CPU/RAM, but shared disks via fast network
- Used in data warehouses
- limited scaling due to locking/conflict management

## Shared nothing architecture

- Horizontal scaling (scale out)
- Each node manages its own CPU/RAM
- Coordination via software/network

## Distribution Goals

- Évolutivité (Scalability) for high volume/load
- Tolérance aux pannes/haute disponibilité for continuous function
- Latence réduite serving users from nearby centers

**Preferred Architecture** Architectures sans partage (Scale Out). Each node manages its own CPU, RAM, and disks. Coordination via network/software. Excellent price/performance.

## Distribution Mechanisms

- Partitionnement (Sharding)** divides large DB into smaller partitions.

- RéPLICATION** maintains copies for latency, availability, and read throughput.

**Single-Leader Replication** One leader handles all writes and sends changes via logs to followers.

**Replication Timing Trade-offs** **Synchronie** guarantees up-to-date followers but adds latency/risk of leader waiting. **Asynchronie** reduces latency but risks data loss if leader fails before replication.

**Leader Failure** Requires **Failover**. A follower is promoted. **Consensus algorithm** (Raft, Paxos) selects the new leader, usually the one with the most recent data.

**Replication Lag/Consistency** Replication delay causes **cohérence éventuelle (eventual consistency)**. Followers eventually catch up if writes stop.

## Read Guarantees

- Reading Your Own Writes:** Requires ensuring the user reads from the leader, or tracking the user's latest update timestamp to guarantee the replica reflects that change.
- Lectures Monotones (Monotonic Reads):** A user never sees an older value after a newer one. Solution: Hash the user ID to ensure reads stick to the same replica.

**Leaderless Replication** Any replica accepts writes. Consistency handled via **Quorums**.

**Quorum Requirement** Need  $W + R > N$  (Writes + Reads > Replicas). This ensures at least one node queried for a read has the latest written data.

**Partitioning by Key Range** Assigns continuous key ranges. Good for **range queries**. Risk of **hot spots** if sequential keys (like timestamps) are frequently written.

**Partitioning by Key Hash** Uses hashing for even key distribution. Eliminates hot spots. Sacrifices efficient range queries.

**Rebalancing** Process of moving load/data between nodes.

- Avoid **hash mod N** because changing  $N$  requires moving most keys.
- Preferred Strategy:** Use a **fixed number of partitions** (more than nodes,  $N$ ). Only entire partitions are moved during rebalancing.

## Concurrency Conflicts (Centralized DB)

- Dirty Write (Écriture Sale):** One transaction overwrites another's uncommitted value. Avoided by acquiring **write locks**.
- Dirty Read (Lecture Sale):** Reading data written by an uncommitted transaction.

**CAP Theorem** System cannot simultaneously guarantee all three: Consistency, Availability, Partition Tolerance.

**CAP Trade-off** Partition Tolerance (P) is unavoidable. When network partitioning occurs, designers must choose between **Consistency** (sacrificing A by rejecting

ing requests) or **Availability** (responding, even if data is stale).